

# 计算机体系结构实验 4 - 5 实验报告

## 实验 4

### 实验流程

### 完善代码

得到以下代码：



```
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int width)
{
    // Calculate the row index of the P element and M int
    row = blockIdx.y * blockDim.y + threadIdx.y;

    // Calculate the column index of the P element and N int
    col = blockIdx.x * blockDim.x + threadIdx.x;

    // Ensure the thread is within bounds if
    (row < width && col < width) {
        float pvalue = 0.0;

        // Each thread computes one element of the matrix
        for(int k = 0; k < width; ++k) {
            pvalue += d_M[row * width + k] * d_N[k * width + col];
        }

        // Store the computed value into the output matrix
        d_P[row * width + col] = pvalue;
    }
}
```



```

for (int j = 0; j < nIter; j++) {
    // matrixMulCPU(reference, h_M, h_N, m, k, n);
    MatrixMulKernel<<<grid, block>>>(d_M, d_N, d_P, m);
    // MatrixMulSharedMemKernel<<<grid, block>>>(d_M, d_N, d_P, m, n);
    // cublasSgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N, n, m, k, &alpha, d_N, n, d_M,
    k, &beta, d_P, n);
}

```

观察发现计算结果正确

```

./MatrixMulKernel 1 1000
Kernel Elapsed Time: 0.557 ms
Performance= 3588.04 GFlop/s, Time= 0.557 msec, Size= 2000000000 Ops
Computing result using host CPU...done.
Listing first 100 Differences > 0.000010...

```

更改矩阵尺寸，对比不同参数下的计算结果

```

(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-mair
./MatrixMulKernel 0 1000
Kernel Elapsed Time: 0.555 ms
Performance= 3603.97 GFlop/s, Time= 0.555 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-mair
./MatrixMulKernel 0 3000
Kernel Elapsed Time: 21.366 ms
Performance= 2527.34 GFlop/s, Time= 21.366 msec, Size= 54000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-mair
./MatrixMulKernel 0 5000
Kernel Elapsed Time: 101.788 ms
Performance= 2456.09 GFlop/s, Time= 101.788 msec, Size= 250000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-mair
./MatrixMulKernel 0 10000
Kernel Elapsed Time: 856.724 ms
Performance= 2334.47 GFlop/s, Time= 856.724 msec, Size= 2000000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-mair
./MatrixMulKernel 0 20000
Kernel Elapsed Time: 8493.021 ms
Performance= 1883.90 GFlop/s, Time= 8493.021 msec, Size= 16000000000000 Ops

```

随着矩阵尺寸的增大，performance 下降

更改tile\_size

```

(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-202
./2 0 1000
Kernel Elapsed Time: 4.222 ms
Performance= 473.71 GFlop/s, Time= 4.222 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-202
./4 0 1000
Kernel Elapsed Time: 1.334 ms
Performance= 1498.93 GFlop/s, Time= 1.334 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-202
./8 0 1000
Kernel Elapsed Time: 0.681 ms
Performance= 2935.57 GFlop/s, Time= 0.681 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-202
./16 0 1000
Kernel Elapsed Time: 0.555 ms
Performance= 3602.80 GFlop/s, Time= 0.555 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-202
./32 0 1000
Kernel Elapsed Time: 0.591 ms
Performance= 3384.09 GFlop/s, Time= 0.591 msec, Size= 2000000000 Ops

```

随着tile\_size增大, performance 升高, 但之后会下降, 但在tile\_size 大于64的时候, 计算开始出错, 如下图:



```

./64 1 1000
Kernel Elapsed Time: 0.001 ms
Performance= 2840909.00 GFlop/s, Time= 0.001 msec, Size= 2000000000 Ops
Computing result using host CPU...done.
Listing first 100 Differences > 0.000010...
  Loc(0,0)    CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(1,0)    CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(2,0)    CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(3,0)    CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(4,0)    CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(5,0)    CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(6,0)    CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(7,0)    CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(8,0)    CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(9,0)    CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(10,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(11,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(12,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(13,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(14,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(15,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(16,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(17,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(18,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(19,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(20,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(21,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(22,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(23,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(24,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(25,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(26,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(27,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(28,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000
  Loc(29,0)   CPU=750.00000    GPU=0.00000    Diff=750.000000
  Loc(30,0)   CPU=375.00000    GPU=0.00000    Diff=375.000000

```

我们可以得到结论：随着矩阵尺寸的增大，performance 下降，随着tile\_size增大，performance 升高，但之后会下降，但在tile\_size 大于64的时候，计算开始出错

## 实验 5

### 实验流程

#### 完善代码

得到以下代码：

```

const int BLOCK_SIZE = TILE_WIDTH;
__global__ void MatrixMulSharedMemKernel(float *A,
    float *B, float *C, int wA,

```

```
    int wB) {  
    // Block index  
    int bx = blockIdx.x;  
    int by = blockIdx.y;  
  
    // Thread index  
    int tx = threadIdx.x;  
    int ty = threadIdx.y;  
  
    // Index of the first sub-matrix of A processed by the block  
    int aBegin = wA * BLOCK_SIZE * by;  
  
    // Index of the last sub-matrix of A processed by the block  
    int aEnd   = aBegin + wA - 1;  
  
    // Step size used to iterate through the sub-matrices of A int  
    aStep = BLOCK_SIZE;  
  
    // Index of the first sub-matrix of B processed by the block  
    int bBegin = BLOCK_SIZE * bx;
```

```

// Step size used to iterate through the sub-matrices of B int bStep = BLOCK_SIZE *
wB;

// Csub is used to store the element of the block sub-matrix
// that is computed by the thread
float Csub = 0;

// Loop over all the sub-matrices of A and B
// required to compute the block sub-matrix
for (int a = aBegin, b = bBegin;
    a <= aEnd; // Ensure all tiles are covered
    a += aStep, b += bStep) {
    // Declaration of the shared memory array As used to
    // store the sub-matrix of A
    __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];

    // Declaration of the shared memory array Bs used to
    // store the sub-matrix of B
    __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];

    // Load the matrices from device memory
    // to shared memory; each **thread** loads
    // one element of each matrix
    // --- TO DO :Load the elements of the sub-matrix of A into As ---
    int aRow = a / wA + ty; // Calculate row in A
    int aCol = a % wA + tx; // Calculate column in A
    if (aRow < wA && aCol < wA)
        As[ty][tx] = A[aRow * wA + aCol];
    else
        As[ty][tx] = 0.0f;

    // --- Load the elements of the sub-matrix of B into Bs ---
    int bRow = b / wB + ty; // Calculate row in B
    int bCol = b % wB + tx; // Calculate column in B
    if (bRow < wA && bCol < wB)
        Bs[ty][tx] = B[bRow * wB + bCol];
    else
        Bs[ty][tx] = 0.0f;

    // synchronize to make sure the matrices are loaded
    __syncthreads();

    // Multiply the two matrices together;
    // each thread computes one element
    // of the block sub-matrix
#pragma unroll
    // --- TO DO :Implement the matrix multiplication using the sub-matrices As
    and Bs ---
    for (int k = 0; k < BLOCK_SIZE; ++k) {
        Csub += As[ty][k] * Bs[k][tx];
    }

    // Synchronize to make sure that the preceding
    // computation is done before loading two new
    // sub-matrices of A and B in the next iteration

```

```

__syncthreads();
}

// write the block sub-matrix to device memory;
// each thread writes one element
int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;
// --- TO DO :Store the computed Csub result into matrix C --- int row_C = by *
BLOCK_SIZE + ty;
int col_C = bx * BLOCK_SIZE + tx; if (row_C < wA && col_C
< wB)
    C[c + ty * wB + tx] = Csub;

```

修改 main() 使得使用 MatrixMulKernel<<<grid, block>>>(d\_M, d\_N, d\_P, m);

```

for (int j = 0; j < nIter; j++) {
    // matrixMulCPU(reference, h_M, h_N, m, k, n);
    // MatrixMulKernel<<<grid, block>>>(d_M, d_N, d_P, m);
    MatrixMulSharedMemKernel<<<grid, block>>>(d_M, d_N, d_P, m, n);
    // cublasSgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N, n, m, k, &alpha, d_N, n, d_M,
k, &beta, d_P, n);
}

```

观察发现计算结果正确

```

(base) inspur@inspur:/data/inspur/workspace-j0hnnny/HITSZ-Comp-Arch-2024-
./MatrixMulSharedMemKernel 1 1000
Kernel Elapsed Time: 0.483 ms
Performance= 4142.03 GFlop/s, Time= 0.483 msec, Size= 2000000000 Ops
Computing result using host CPU...done.
Listing first 100 Differences > 0.000010...

Total Errors = 0

```

更改矩阵尺寸，对比不同参数下的计算结果



```

(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./MatrixMulSharedMemKernel 0 1000
Kernel Elapsed Time: 0.483 ms
Performance= 4144.29 GFlop/s, Time= 0.483 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./MatrixMulSharedMemKernel 0 3000
Kernel Elapsed Time: 12.652 ms
Performance= 4268.07 GFlop/s, Time= 12.652 msec, Size= 54000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./MatrixMulSharedMemKernel 0 5000
Kernel Elapsed Time: 62.401 ms
Performance= 4006.37 GFlop/s, Time= 62.401 msec, Size= 250000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./MatrixMulSharedMemKernel 0 10000
Kernel Elapsed Time: 497.037 ms
Performance= 4023.84 GFlop/s, Time= 497.037 msec, Size= 2000000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./MatrixMulSharedMemKernel 0 20000
Kernel Elapsed Time: 4010.275 ms
Performance= 3989.75 GFlop/s, Time= 4010.275 msec, Size= 16000000000000 Ops

```

可以看见，随着矩阵尺寸的增大，performance也下降

更改 tile\_size，对比不同参数下的计算结果

```

./2 0 1000
Kernel Elapsed Time: 19.070 ms
Performance= 104.88 GFlop/s, Time= 19.070 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./4 0 1000
Kernel Elapsed Time: 2.557 ms
Performance= 782.25 GFlop/s, Time= 2.557 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./8 0 1000
Kernel Elapsed Time: 0.712 ms
Performance= 2809.92 GFlop/s, Time= 0.712 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./16 0 1000
Kernel Elapsed Time: 0.480 ms
Performance= 4163.45 GFlop/s, Time= 0.480 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main
./32 0 1000
Kernel Elapsed Time: 0.480 ms
Performance= 4164.72 GFlop/s, Time= 0.480 msec, Size= 2000000000 Ops

```

调用自带的矩阵乘法算子，编译运行



```
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 1 1000
Kernel Elapsed Time: 216.314 ms
Performance= 9.25 GFlop/s, Time= 216.314 msec, Size= 2000000000 Ops
Computing result using host CPU...done.
Listing first 100 Differences > 0.000010...

Total Errors = 0
```

更改矩阵尺寸，对比不同参数下的计算结果

```
./cublasSgemm 0 1000
Kernel Elapsed Time: 204.143 ms
Performance= 9.80 GFlop/s, Time= 204.143 msec, Size= 2000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 0 3000
Kernel Elapsed Time: 208.129 ms
Performance= 259.46 GFlop/s, Time= 208.129 msec, Size= 54000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 0 5000
Kernel Elapsed Time: 219.477 ms
Performance= 1139.07 GFlop/s, Time= 219.477 msec, Size= 250000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 0 10000
Kernel Elapsed Time: 298.059 ms
Performance= 6710.07 GFlop/s, Time= 298.059 msec, Size= 2000000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 0 20000
Kernel Elapsed Time: 1187.059 ms
Performance= 13478.69 GFlop/s, Time= 1187.059 msec, Size= 16000000000000 Ops
(base) inspur@inspur:/data/inspur/workspace-johnny/HITSZ-Comp-Arch-2024-main/1
./cublasSgemm 0 40000
Kernel Elapsed Time: 8640.741 ms
Performance= 14813.54 GFlop/s, Time= 8640.741 msec, Size= 128000000000000 Ops
```

Tile\_width不影响cublasSgemm，故不进行测试。

## 实验结果及原理分析

方法	矩阵尺寸	TILE_WIDTH	Performance(GFlop/s)
MatrixMulKernel	1000	2	473
MatrixMulKernel	1000	4	1499
MatrixMulKernel	1000	8	2936
MatrixMulKernel	1000	16	3602
MatrixMulKernel	1000	32	3384
MatrixMulKernel	3000	32	2527
MatrixMulKernel	5000	32	2456
方法	矩阵尺寸	TILE_WIDTH	Performance(GFlop/s)

MatrixMulKernel	10000	32	2334
MatrixMulKernel	20000	32	1883
MatrixMulSharedMemKernel	1000	2	108
MatrixMulSharedMemKernel	1000	4	782
MatrixMulSharedMemKernel	1000	8	2809
MatrixMulSharedMemKernel	1000	16	4163
MatrixMulSharedMemKernel	1000	32	4144
MatrixMulSharedMemKernel	3000	32	4268
MatrixMulSharedMemKernel	5000	32	4006
MatrixMulSharedMemKernel	10000	32	4023
MatrixMulSharedMemKernel	20000	32	4010
cublasSgemm	500		1.3
cublasSgemm	1000		9.8
cublasSgemm	3000		259
cublasSgemm	5000		1139
cublasSgemm	10000		6710
cublasSgemm	20000		13478
cublasSgemm	40000		14813

分析数据我们可以知道MatrixMulSharedMemKernel比 MatrixMulKernel表现要好，而且随着举证变大性能下降并不明显cublasSgemm在矩阵较小的时候性能并不好，比其他两种方法相差很多，但是随着矩阵增大，其性能也升高明显，最后比其他两种方法优秀很多。

分析原理

MatrixMulSharedMemKernel有效地减少了全局内存访问次数，而共享内存的性能比全局内存强很多，并且更好地利用了GPU的缓存层次结构，因此性能更好。

cublasSgemm在小矩阵时性能不佳是因为其优化针对大规模并行计算，存在启动开销；随着矩阵增大，其高效利用GPU资源和高度优化的算法特性得以充分发挥，从而展现出显著的性能优势。